


Research is the part that needs a human. *Everything else is mechanics.*

CRA runs the full arc of clinical research — evidence synthesis, trial design, regulatory start-up, data collection, statistical analysis, and publication drafting — on one governed platform, under the standards a submission must survive: ICH GCP, 21 CFR Part 11, PRISMA 2020, CDISC. Your researchers put their judgment on the science. The platform does the rest, and proves it did it correctly.



END TO END

One platform, the whole research lifecycle

Fifteen specialist AI workflows hand off to each other with full context — the meta-analysis feeds the GRADE tables, the SAP feeds the study report, the locked database feeds the statistics.

I Evidence Literature search strategy, systematic-review protocol, dual-reviewer screening, meta-analysis, network & IPD meta-analysis, risk of bias, GRADE, living reviews.	II Design PICOT framing, sample size & power, full Statistical Analysis Plan drafted to ICH E9.
III Start-up Registry records (ClinicalTrials.gov, EU CTIS), IRB packet & consent forms checked against 21 CFR §50.25, randomisation schedules.	IV Execution EDC + ePRO capture, AE/SAE with the 24-hour clock, drug accountability, lab feeds, visit calendars, monitoring & multi-site rollups — Part 11 signatures and audit throughout.
V Reporting Post-lock statistics, CDISC SDTM/ADaM datasets, ICH E3 study report, journal manuscript, plain-language summaries.	

15	1,650+	42	7+2	24 h
specialist workflows	automated tests	PRISMA 2020 items enforced	SDTM domains + ADaM datasets	SAE regulatory clock

METHODOLOGICAL INTEGRITY

This is not “AI-generated research”

The distinction that matters to a reviewer: who did the science, and can every number be traced. CRA is built so the answers are *your researchers* and *yes*.

The AI automates the mechanical work and structures the workflow for compliance. It is architecturally prevented from inventing the science.

Citations only from real searches

The platform rejects any answer quoting effect sizes, confidence intervals, or PMIDs from the model's memory. Every citation traces to an actual literature search executed during the session.

Statistics computed, never estimated

Every pooled effect, survival curve, and forest plot is produced by real statistical code running on real data in an isolated sandbox — the model orchestrates the analysis; it cannot make up the result.

Provenance on every number

Each figure in a study report carries a machine-checked reference to the source artefact it was derived from. GRADE certainty is computed from the domain ratings — never asserted.

Structured, stage-gated workflows

A meta-analysis cannot skip from question to conclusion: the protocol, search, screening, and extraction stages are enforced in order, with the researcher confirming each gate.

21 CFR Part 11

ICH E6 (GCP)

ICH E2A

ICH E3

ICH E9

ALCOA+

PRISMA 2020

GRADE

CDISC SDTM / ADaM / Define-XML

CONSORT

IQ / OQ / PQ validation pack

WHERE THE TIME GOES

Offload the mechanics, keep the thinking

Most of a study's calendar time is consumed by work that is essential, skilled — and mechanical. CRA absorbs it.

Hours, not weeks

Boolean search strategies with MeSH expansion, PRISMA-P protocols, screening logs, extraction tables, SAP boilerplate, registry field mapping, consent-form compliance checks — drafted in minutes, reviewed by your team, finalized the same day.

Compliance as a property, not a project

Audit trails, electronic signatures with password re-authentication, delegation logs, separation-of-duties checks, and inspector-ready PDF reports accumulate automatically as the team works — not as a pre-submission scramble.

Nothing falls through

The serious-adverse-event clock, visit windows, query resolution, drug reconciliation, and training expiries are tracked by the system and surfaced before they become findings.

Higher-order work stays human

Framing the question, judging clinical relevance, interpreting heterogeneity, deciding what the evidence means — the work that makes research *research* gets the researcher's full attention.

COST CONTROL

Budgets you set. Spend you can see.

CRA runs on Claude models via Amazon Bedrock inside your own AWS account. You choose the model tier per deployment, cap the spend, and watch per-workflow cost rollups in real dollars.

Claude Haiku

Small team

~\$500 / month

A lean research group runs a substantial program — literature reviews, meta-analyses, protocol and SAP drafting — on the fastest, most economical tier.

Claude Sonnet

Department

Scaled to workload

Deeper reasoning for complex synthesis, multi-arm network meta-analysis, and higher concurrent usage across a department.

Claude Opus

Enterprise

Premium tier

Maximum-capability model for the hardest synthesis and reporting work, with the same guardrails and audit posture.

Daily token quotas, per-trial dollar budgets, and per-workflow / per-model cost attribution are built in — spend stops at the cap you set, and every dollar is traceable to the work it produced.

EVIDENCE ACCESS & OPERATIONS

The literature your license already covers

Open + subscription sources

PubMed and Europe PMC are built in. Connect your institution’s publisher credentials — Elsevier titles such as *The Lancet*, and other subscription sources — with per-source API keys, rate limits, and usage caps managed from the admin console.

Your private evidence library

Upload PDFs and cached publications into a local, searchable library; the AI grounds its answers in what your team has actually collected, with citations back to each passage.

Your cloud, your data

Deploys into your own AWS account — private network, single sign-on, role-based access with separation-of-duties, and PHI kept in a dedicated store. Nothing leaves your environment.

Every deliverable, export-ready

Protocols, SAPs, study reports, GRADE tables, manuscripts, and regulatory logs export as formatted PDF and Word documents — submission-shaped from day one.

CAPABILITY INDEX

What ships in the platform today

EVIDENCE SYNTHESIS

Search strategy & MeSH expansion
SR protocol (PRISMA-P / PROSPERO)
Dual-reviewer screening + adjudication
Pairwise, network & IPD meta-analysis
Risk of bias (RoB 2, ROBINS-I)
GRADE summary-of-findings
Living reviews with panel voting

DESIGN & START-UP

Sample size & power (4 designs)
Statistical Analysis Plan drafter
CT.gov + EU CTIS registration records
IRB packet & multilingual consent forms
Randomisation & IRT (4 algorithms, blinding, code-break)

CONDUCT (EDC)

Form builder, edit checks, queries, SDV
ePRO via participant magic links
AE/SAE (ICH E2A) + deviations/CAPA
Recruitment log with CONSORT funnel
Visit calendars + email reminders
Drug accountability & reconciliation
Lab feeds: HL7 v2, CDISC LAB, FHIR R4
Multi-site KPI rollups

ANALYSIS & REPORTING

Post-lock statistics: K-M, Cox, MMRM, subgroups
CDISC SDTM / ADaM / Define-XML bundle
ICH E3 clinical study report
Journal manuscript + reviewer responses
Plain-language summaries (readability-scored)
Funnel, waterfall, swimmer & forest plots

GOVERNANCE

SSO (Amazon Cognito), RBAC, SoD matrix
Part 11 signatures + append-only audit
Delegation log with PI countersigning
Onboarding gates on clinical writes
IQ/OQ/PQ validation pack + traceability
Portfolio & budget dashboards

The credibility of your research is the product.

CRA exists so that when your team submits, the methodology holds — every search reproducible, every number traceable, every signature compliant. Faster because the mechanics are automated; publishable because the science never was.

CRA · Clinical Research Assistant — Executive brief

Companion documents: Feature Guide · Demo Guide