


Research is the part that needs a human.
Everything else is mechanics.

CRA runs the full arc of clinical research — evidence synthesis, trial design, regulatory start-up, data collection, statistical analysis, and publication drafting — on one governed platform, aligned with the standards reviewers and regulators expect: ICH GCP, 21 CFR Part 11, PRISMA 2020, CDISC. Your researchers put their judgment on the science. The platform does the rest, and keeps the auditable record of how it was done.



END TO END

One platform, the whole research lifecycle

Fifteen specialist AI workflows hand off to each other with full context — the meta-analysis feeds the GRADE tables, the SAP feeds the study report, the locked database feeds the statistics.

I Evidence Literature search strategy, systematic-review protocol, dual-reviewer screening, meta-analysis, network & IPD meta-analysis, risk of bias, GRADE, living reviews.	II Design PICOT framing, sample size & power, full Statistical Analysis Plan drafted to ICH E9.
III Start-up Registry records (ClinicalTrials.gov, EU CTIS), IRB packet & consent forms checked against 21 CFR §50.25, randomisation schedules.	IV Execution EDC + ePRO capture, AE/SAE with the 24-hour clock, drug accountability, lab feeds, visit calendars, monitoring & multi-site rollups — Part 11 signatures and audit throughout.
V Reporting Post-lock statistics, CDISC SDTM/ADaM datasets, ICH E3 study report, journal manuscript, plain-language summaries.	

15	1,650+	42	7+2	24 h
specialist workflows	automated tests	PRISMA 2020 items enforced	SDTM domains + ADaM datasets	SAE regulatory clock

METHODOLOGICAL INTEGRITY

This is not “AI-generated research”

The distinction that matters to a reviewer: who did the science, and can every number be traced. CRA is built so the answers are *your researchers* and *yes*.

The AI automates the mechanical work and structures the workflow for compliance. It is architecturally prevented from inventing the science.

Citations only from real searches

The platform rejects any answer quoting effect sizes, confidence intervals, or PMIDs from the model's memory. Every citation traces to an actual literature search executed during the session.

Statistics computed, never estimated

Every pooled effect, survival curve, and forest plot is produced by real statistical code running on real data in an isolated sandbox — the model orchestrates the analysis; it cannot make up the result.

Provenance on every number

Each figure in a study report carries a machine-checked reference to the source artefact it was derived from. GRADE certainty is computed from the domain ratings — never asserted.

Structured, stage-gated workflows

A meta-analysis cannot skip from question to conclusion: the protocol, search, screening, and extraction stages are enforced in order, with the researcher confirming each gate.

21 CFR Part 11

ICH E6 (GCP)

ICH E2A

ICH E3

ICH E9

ALCOA+

PRISMA 2020

GRADE

CDISC SDTM / ADaM / Define-XML

CONSORT

IQ / OQ / PQ validation pack

WHERE THE TIME GOES

Offload the mechanics, keep the thinking

Most of a study's calendar time is consumed by work that is essential, skilled — and mechanical. CRA absorbs it.

Hours, not weeks

Boolean search strategies with MeSH expansion, PRISMA-P protocols, screening logs, extraction tables, SAP boilerplate, registry field mapping, consent-form compliance checks — drafted in minutes, reviewed by your team, finalized the same day.

Compliance as a property, not a project

Audit trails, electronic signatures with password re-authentication, delegation logs, separation-of-duties checks, and inspector-ready PDF reports accumulate automatically as the team works — not as a pre-submission scramble.

Nothing falls through

The serious-adverse-event clock, visit windows, query resolution, drug reconciliation, and training expiries are tracked by the system and surfaced before they become findings.

Higher-order work stays human

Framing the question, judging clinical relevance, interpreting heterogeneity, deciding what the evidence means — the work that makes research *research* gets the researcher's full attention.

COST CONTROL

Budgets you set. Spend you can see.

CRA runs on Claude models via Amazon Bedrock inside your own AWS account. You choose the model tier per deployment, cap the spend, and watch per-workflow cost rollups in real dollars.

Claude Haiku

Small team

~\$500 / month

A lean research group runs a substantial program — literature reviews, meta-analyses, protocol and SAP drafting — on the fastest, most economical tier.

Claude Sonnet

Department

Scaled to workload

Deeper reasoning for complex synthesis, multi-arm network meta-analysis, and higher concurrent usage across a department.

Claude Opus

Enterprise

Premium tier

Maximum-capability model for the hardest synthesis and reporting work, with the same guardrails and audit posture.

Daily token quotas, per-trial dollar budgets, and per-workflow / per-model cost attribution are built in — spend stops at the cap you set, and every dollar is traceable to the work it produced.

EVIDENCE ACCESS & OPERATIONS

The literature your license already covers

Open + subscription sources

PubMed and Europe PMC are built in. Connect your institution’s publisher credentials — Elsevier titles such as *The Lancet*, and other subscription sources — with per-source API keys, rate limits, and usage caps managed from the admin console.

Your private evidence library

Upload PDFs and cached publications into a local, searchable library; the AI grounds its answers in what your team has actually collected, with citations back to each passage.

Your cloud, your data

Deploys into your own AWS account — private network, single sign-on, role-based access with separation-of-duties, and PHI kept in a dedicated store. Model inference runs on Amazon Bedrock inside your account — research data never transits a third-party AI service.

Every deliverable, export-ready

Protocols, SAPs, study reports, GRADE tables, manuscripts, and regulatory logs export as formatted PDF and Word documents — submission-shaped from day one.

Claims you can ask us to show you

Nothing on these pages is aspirational copy. Each commitment below maps to an artefact — a validation document, an audit trail, a sign-off gate — that we will walk through with your review committee.

The researcher remains the author

Every workflow pauses at human gates. The statistical analysis plan and post-lock analyses are confirmed by your statistician; AE/SAE classifications are proposed by the system and confirmed by the investigator; registry, IRB, and consent documents are drafts for your regulatory reviewer; CDISC datasets are released by your data manager. The AI never signs anything — people do, with Part 11 signatures.

Validated like the systems it replaces

The platform ships with an IQ/OQ/PQ validation pack and a requirements-traceability matrix, backed by 1,650+ automated tests run on every change. GRADE certainty, readability scores, and all statistics are computed by verifiable code — never asserted by the model.

Security & privacy by design

Single sign-on, role-based access with separation-of-duties (12 roles, over a hundred distinct permissions), PHI in a dedicated store, append-only audit trails, and model-authored code confined to an isolated, network-disabled sandbox — all inside your own AWS account.

Fits the systems you already run

Lab data in via HL7 v2, FHIR R4, and CDISC LAB; source documents via CSV with checksum de-duplication; references via BibTeX/RIS; deliverables out as PDF, Word, and Define-XML. Deployment is infrastructure-as-code into your AWS environment.

Operational metrics — hours saved per protocol, screening throughput, query cycle times — will come from structured pilot evaluations and be reported as measured, not projected. If your group would like to be among the first pilot sites, we would welcome the scrutiny.

What ships in the platform today

EVIDENCE SYNTHESIS

Search strategy & MeSH expansion
SR protocol (PRISMA-P / PROSPERO)
Dual-reviewer screening + adjudication
Pairwise, network & IPD meta-analysis
Risk of bias (RoB 2, ROBINS-I)
GRADE summary-of-findings
Living reviews with panel voting

DESIGN & START-UP

Sample size & power (4 designs)
Statistical Analysis Plan drafter
CT.gov + EU CTIS registration records
IRB packet & multilingual consent forms
Randomisation & IRT (4 algorithms, blinding, code-break)

CONDUCT (EDC)

Form builder, edit checks, queries, SDV
ePRO via participant magic links
AE/SAE (ICH E2A) + deviations/CAPA
Recruitment log with CONSORT funnel
Visit calendars + email reminders
Drug accountability & reconciliation
Lab feeds: HL7 v2, CDISC LAB, FHIR R4
Multi-site KPI rollups

ANALYSIS & REPORTING

Post-lock statistics: K-M, Cox, MMRM, subgroups
CDISC SDTM / ADaM / Define-XML bundle
ICH E3 clinical study report
Journal manuscript + reviewer responses
Plain-language summaries (readability-scored)
Funnel, waterfall, swimmer & forest plots

GOVERNANCE

SSO (Amazon Cognito), RBAC, SoD matrix
Part 11 signatures + append-only audit
Delegation log with PI countersigning
Onboarding gates on clinical writes
IQ/OQ/PQ validation pack + traceability
Portfolio & budget dashboards

The credibility of your research is the product.

CRA is engineered so that when your team submits, the methodology stands up to review — searches built to be reproducible, numbers that carry their provenance, signatures that meet Part 11. Faster because the mechanics are automated; credible because the science never left your researchers' hands.